

§ MongoDB to PostgreSQL Migration Plan

§ Phase I: Analysis and Schema Design

This is the most critical phase. A thorough analysis and a well-designed relational schema will prevent significant issues later.

1. Analyze MongoDB Schema:

- Document all MongoDB collections and the structure of the documents within them.
- Identify embedded documents and arrays. These will need to be modeled as separate tables or using specific PostgreSQL types (like `JSONB` or `ARRAY`).
- Map out relationships between collections (e.g., manual references using `ObjectId`).

2. Design Target PostgreSQL Schema:

- For each MongoDB collection, design one or more corresponding PostgreSQL tables.
- Define columns with appropriate PostgreSQL data types. Create a mapping table.

MongoDB Type	PostgreSQL Type	Notes
<code>ObjectId</code>	<code>UUID</code> or <code>BIGSERIAL</code>	<code>UUID</code> is a good choice to maintain global uniqueness. <code>BIGSERIAL</code> is simpler for auto-incrementing primary keys.
<code>String</code>	<code>TEXT</code> or <code>VARCHAR(n)</code>	Use <code>TEXT</code> unless you have a strict length requirement.
<code>Number</code>	<code>INTEGER</code> , <code>BIGINT</code> , <code>NUMERIC</code>	Choose based on the range and precision of the numbers.
<code>Date</code>	<code>TIMESTAMP WITH TIME ZONE</code>	<code>TIMESTAMPTZ</code> is best practice for storing time data.
<code>Boolean</code>	<code>BOOLEAN</code>	Direct mapping.
<code>Array</code>	<code>ARRAY</code> type (e.g., <code>TEXT[]</code>) or a separate join table.	A join table is more flexible and normalized for complex objects.
<code>Object</code> / <code>Embedded</code>	<code>JSONB</code> or a separate table with a foreign key	<code>JSONB</code> is powerful for unstructured data, but a separate table is better for

MongoDB Type	PostgreSQL Type	Notes
Doc	relationship.	data you need to query or index directly.

3. Define Relationships:

- Convert MongoDB document references into PostgreSQL foreign key constraints to enforce relational integrity.
- Use join tables to model many-to-many relationships.

4. Choose Migration Tools:

- Decide between using a dedicated ETL tool (e.g., Pentaho, Talend) or writing custom migration scripts (e.g., in Python with `pymongo` and `psycopg2`, or Node.js with `mongodb` and `pg`). Custom scripts offer more control, which is often necessary for complex transformations.

§ Phase 2: Environment Setup and Schema Implementation

1. **Set up PostgreSQL:** Install and configure a PostgreSQL server for development and testing. Create a new database, users, and roles with appropriate permissions.
2. **Write DDL Scripts:** Based on the design from Phase 1, write the SQL Data Definition Language (DDL) scripts.
3. **Execute DDL:** Run the scripts to create all tables, indexes, sequences, views, and foreign key constraints in the target PostgreSQL database.

```
-- Example DDL for migrating a 'users' collection

CREATE EXTENSION IF NOT EXISTS "uuid-oss"; -- To generate UUIDs

CREATE TABLE users (
    id UUID PRIMARY KEY DEFAULT uuid_generate_v4(),
    email TEXT NOT NULL UNIQUE,
    password_hash TEXT NOT NULL,
    first_name VARCHAR(255),
    last_name VARCHAR(255),
    created_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),
    updated_at TIMESTAMPTZ NOT NULL DEFAULT NOW()
);

-- Example for a related 'posts' collection
CREATE TABLE posts (
```

```
id BIGSERIAL PRIMARY KEY,  
author_id UUID NOT NULL REFERENCES users(id) ON DELETE CASCADE,  
title TEXT NOT NULL,  
content TEXT,  
published_at TIMESTAMPTZ,  
created_at TIMESTAMPTZ NOT NULL DEFAULT NOW(),  
updated_at TIMESTAMPTZ NOT NULL DEFAULT NOW()  
);  
  
CREATE INDEX idx_posts_author_id ON posts(author_id);
```

§ Phase 3: Data Migration (ETL Process)

This phase involves writing and running scripts to Extract, Transform, and Load data.

1. **Extract:** Write code to connect to the MongoDB database and read all documents from a source collection.
2. **Transform:** For each MongoDB document, write transformation logic to:
 - Flatten the document structure to fit the relational model.
 - Convert data types (e.g., `ObjectId` to a string representation for `UUID` or mapping to a new `SERIAL` key).
 - Handle nested arrays and objects by preparing data for insertion into related tables.
3. **Load:** Write code to connect to PostgreSQL and insert the transformed data into the appropriate tables.
 - Be mindful of the insertion order to respect foreign key constraints (e.g., insert `users` before their `posts`).
 - Use transactions to ensure data is inserted atomically.

§ Phase 4: Application Code Refactoring

This is a significant development effort that can be done in parallel with data migration scripting.

1. **Update Dependencies:** Replace the MongoDB driver and ODM/ORM (e.g., Mongoose) with a PostgreSQL driver and ORM (e.g., Sequelize, TypeORM, pg-promise).

2. **Rewrite Data Access Layer:** Refactor all database queries from the MongoDB query syntax to SQL. This includes simple CRUD operations and complex aggregation pipelines.

- **Before (Mongoose):** `User.find({ age: { $gt: 21 } })`
- **After (Sequelize):** `User.findAll({ where: { age: { [Op.gt]: 21 } } })`
- **After (Raw SQL):** `SELECT * FROM users WHERE age > 21`

3. **Adapt Business Logic:** Adjust any application logic that relied on MongoDB-specific features or schemaless flexibility.

§ Phase 5: Validation and Testing

1. Data Integrity Verification:

- Compare row counts between MongoDB collections and PostgreSQL tables.
- Write scripts to perform spot checks, comparing a sample of records from the source and target to ensure data was not corrupted.
- Validate that all relations were created correctly.

2. Application Testing:

- **Unit Tests:** Update and run all unit tests for the refactored data access layer.
 - **Integration Tests:** Perform end-to-end testing of all application features to ensure they work correctly with PostgreSQL.
 - **Performance Testing:** Benchmark the application against the new database. Identify and optimize slow-running SQL queries using tools like `EXPLAIN ANALYZE`. Add indexes where necessary.
-

§ Phase 6: Deployment and Cutover

1. Plan the Cutover:

- **Strategy:** Decide on a "big bang" (downtime required) or a "live" migration strategy. For most applications, a scheduled downtime is safer and simpler.
- **Timeline:** Schedule the migration during a low-traffic period.

2. Execute the Final Migration:

- Put the application into maintenance mode.
- Perform a final, complete data migration from the production MongoDB to the production PostgreSQL database to ensure data is up-to-date.

- Run verification scripts to confirm the integrity of the production data.

3. **Go Live:**

- Update the application's environment variables to point to the new PostgreSQL database.
- Deploy the refactored application code.
- Disable maintenance mode and bring the application back online.

4. **Post-Migration:**

- Closely monitor application logs and performance metrics for any issues.
- Keep the MongoDB database online in a read-only state for a short period as a fallback, but do not allow the application to write to it.
- Once confident, decommission the old MongoDB database.